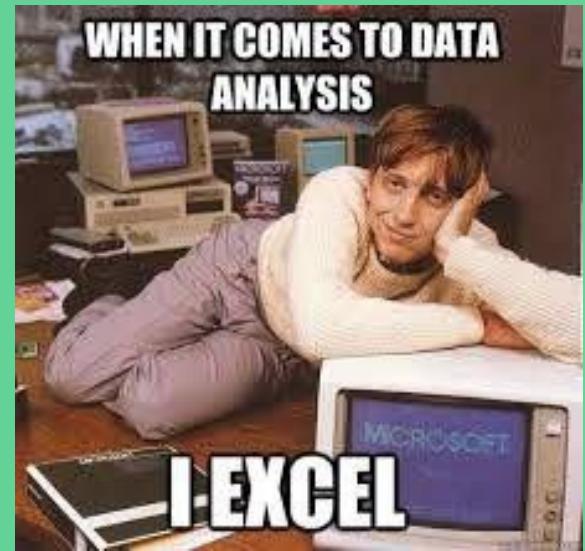


Wireless Analysis with Kismet and ELK

Chase Peterson
@developingchase
github.com/developingchase/wawkelk

I want to understand the wireless environment around me --

- 1) Capture all the RFs with Kismet
- 2) Dive deeper - Spatial, temporal, trend analysis



Overview

Goal: Present lessons learned and a quick start guide to help others and generate interest and involvement

1. Kismet Collection
2. Elasticsearch
3. Logstash
4. Python
5. Kibana
6. Analysis

Kismet

<https://kismetwireless.net>

<https://kismetwireless.net/docs>

@kismetwireless

Discord: <https://discord.gg/5N4ME9a>

Dragorn's recent SharkFest Presentation:

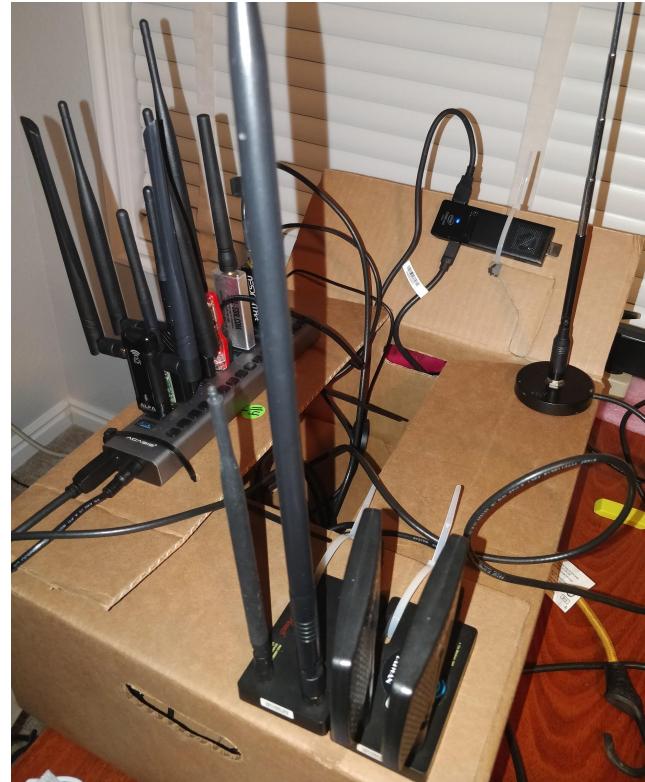
<https://www.youtube.com/watch?v=z6MzIDwjUmc>



Kismet Collection Platforms

“War-Box”/Static Collection Platform

- Intel Compute Stick 325
 - Ubuntu 20.04.01
- 16 Port USB Hub
- Panda PAU09 x 1 (Channel locked)
- Alfa AWUS036ACM x 1 (Channel locked)
- Rosewill x 2 (Scanning)
- Sena UD100
- Ubertooth
- Internal Bluetooth
- NooElec SDR (ADS-B)
- RTL-SDR v3 (RTL-433)
 - rtl-433 dependency
- USB VFAN BG-803 (Ublox) GPS x 1



Kismet Collection Platforms

“War-Pack”

- Raspberry Pi 3
 - Rasbian
 - Re4son Kernel
- Panda PAU09 x 2
- NooElec SDR x 1
- Anker 20000 MHa Battery
- GlobalSat BU-353 S4
- Bagsmart Travel Bag



Kismet Config Tips

Config files - kismet_site.conf

```
server_name=compute_drone
```

```
log_prefix=/home/compute/kismet/
```

```
log_title=homecaptures
```

```
log_types=kismet,pcapppi
```

```
gps=gpsd:host=localhost,port=2947
```

```
source=wlx00c0caab9f88:name=AlfaACM,channels="1,6,11,36,44,149,153,157,161",vht_channels=false,ht_channels=false
```

```
source=rtl433-00009002:name=RTLSDRv3Int
```

```
source_stagger_threshold=6
```

```
hidedata=true
```

```
kis_log_device_filter=IEEE802.11,aa:bb:cc:dd:ee:ff,block
```

```
kis_log_packet_filter=IEEE802.11,aa:bb:cc:dd:ee:ff,block
```

Kismet Data

Logging Data Output Types

Foo.kismetdb - sqitedb; RDB with JSON blob

Foo.pcapppi - PPI has GPS and signal details

Foo.pcapng - Types combined, Radiotap

Data processing tools

kismetdb_clean	kismetdb_dump_devices
kismetdb_statistics	kismetdb_strip_packets
kismetdb_to_gpx	kismetdb_to_kml
kismetdb_to_pcap	kismetdb_to_wiglecsv

API

```
curl http://username:password@localhost:2501/system/status.prettyjson
```

http://localhost:2501/system/tracked_fields.html

Serialization: JSON vs EKJSON

Example Data Fields

[kismet.device.base] - Fields common to all devices

- kismet.device.base.macaddr - mac address
- kismet.device.base.type - printable device type
- kismet.device.base.manuf - manufacturer name

[dot11] - Fields specific to 802.11 devices; lots of nested fields

- dot11.probedssid.ssid - probed SSID
- dot11.device.dot11.device.last.probed.ssid.record.dot11.p
robssid.ssid - Actual nested field result

[rtl433] - Fields for rtl433 devices

- rtl433.device.lightning_strike_count - Strike count

Elasticsearch, Logstash, and Kibana (ELK) Stack

Elasticsearch - Document database built for scaling

Logstash - Data pipeline tool that parses files/logs and pushes them into Elasticsearch

Kibana - Web based visualization tool that displays Elasticsearch data

Windows/Mac/Linux

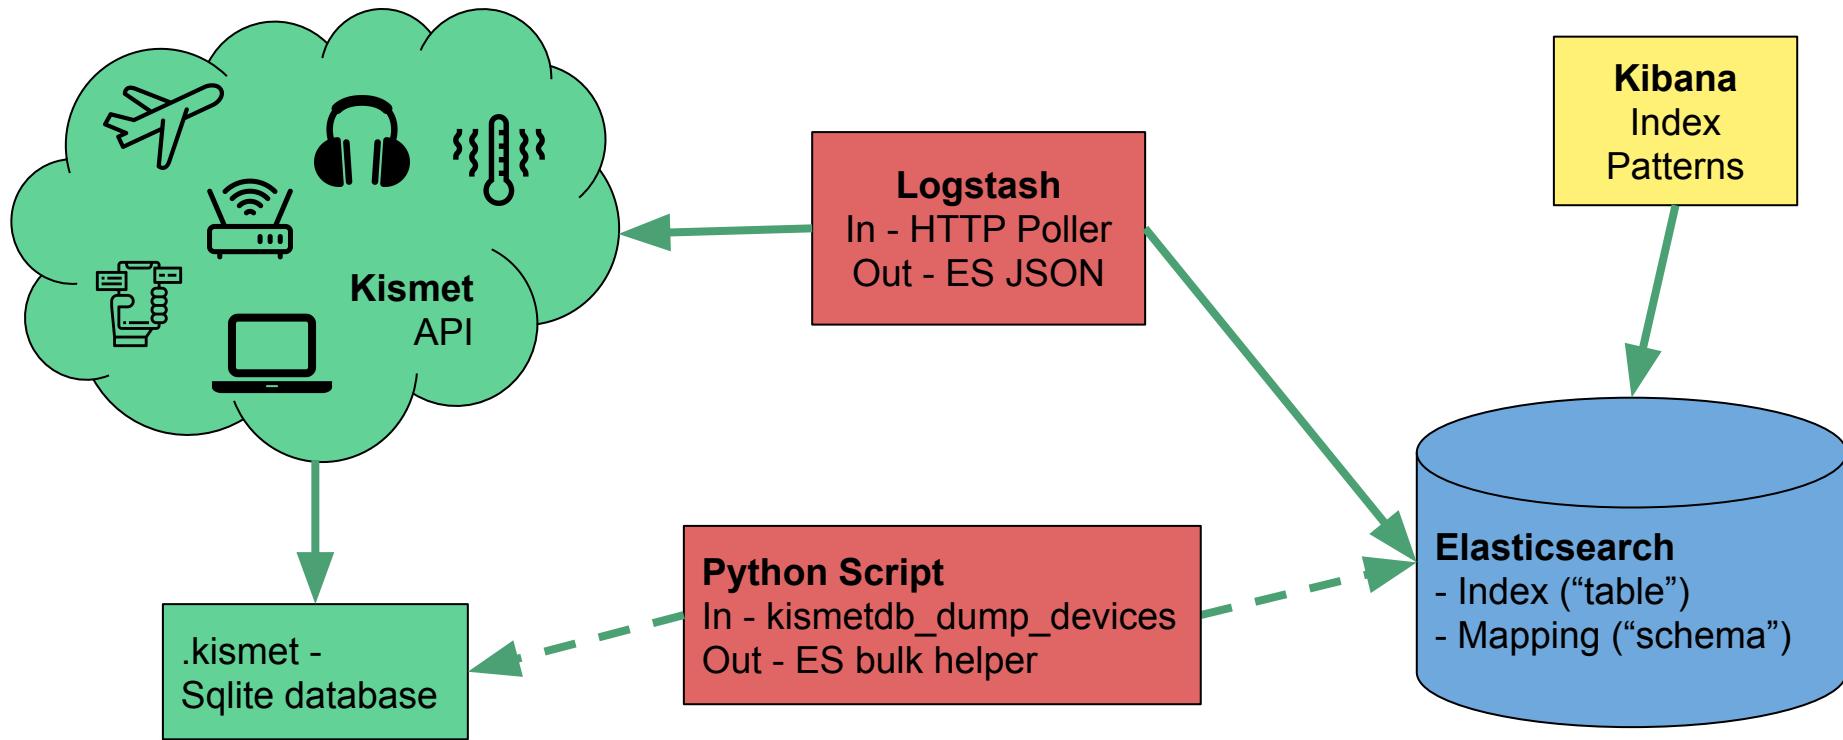
ELK Software Options

- FOSS Package (Limited)
- Basic License (Better)
- Gold, Platinum, Etc...

Deployment Options

- Computer/VM Self Managed
- **Bitnami Virtual Machine Images**
- AWS Cloud
- Elastic.co Cloud (Managed/Self Managed)

WAWKELK Service Diagram



Elasticsearch

Defining Indices

- Lots of optimization settings; great documentation

Defining Mappings

- Static versus Dynamic Mapping
- Must be defined in advance; once data is in the index, it cannot be changed*

Use the Dev Console in Kibana

Test, test, and test again; once happy, blow away everything and start fresh

ES suffers when fields exceed 1000

Nested fields are *terrific*...

Kibana won't detect nested geo_point fields, so you have to pull them into the base document

```
42 PUT kismet_device_activity
43 {
44   "settings": {
45     "index.mapping.ignore_malformed": true
46   },
47   "mappings": {
48     "dynamic": true,
49     "numeric_detection": true,
50     "properties": {
51       "kismet_device_base_location_avg_geopoint": {
52         "type": "geo_point"
53     },
54     "kismet_device_base_location": {
55       "type": "nested",
56       "properties": {
57         "kismet_common_location_avg_loc": {
58           "type": "nested",
59           "properties": {
60             "kismet_common_location_geopoint": {
61               "type": "geo_point"
62             }
63           }
64         }
65       }
66     },
67     "kismet_device_base_tags": {
68       "type": "nested",
69       "properties": {
70         "notes": {
71           "type": "text"
72         }
73       }
74     }
75   }
76 }
77 }
```

Logstash

Runs as a service

Inputs

- Lots of options; HTTP Poller best met this use case

Filters

- Data mangling

Outputs

- Elasticsearch

Kismet-isms

- EKJSON - 1 document per line
- Request fields in “.” notation
- Receive fields in “_” notation

Logstash - Inputs

```
input {  
    http_poller {  
        urls => {  
            device_activity => {  
                url => "http://192.168.0.100:2501/devices/last-time/-305/devices.ekjson"  
                method => post  
                auth => { user => "kismet" password => "password" }  
                body =>  
                    "json=[\"fields\":[\"kismet.device.base.type\",\"kismet.device.base.location\",\"kismet.device.base.tags\",\"kismet.device.base.manuf\",\"kismet.device.base.num_alerts\",\"kismet.device.base.server_uuid\",\"kismet.device.base.last_time\",\"kismet.device.base.frequency\",\"kismet.device.base.channel\",\"kismet.device.base.crypt\",\"kismet.device.base.commonname\",\"kismet.device.base.name\",\"kismet.device.base.phyname\",\"kismet.device.base.macaddr\",\"kismet.device.base.key\"]]  
                headers => {  
                    "Accept" => "application/json"  
                    "Content-Type" => "application/x-www-form-urlencoded"  
                }  
            }  
        }  
    }  
    add_field => { "[@metadata][dest_index]" => "kismet_device_activity" }  
    add_field => { "kismet_host_inline" => "system-one" }  
    add_field => { "kismet_location_inline" => "home" }  
    request_timeout => 8  
    #schedule => { cron => "* * * * * UTC" }  
    schedule => { every => "5m" }  
    codec => "json_lines"  
}  
}
```

1

2

3

4

5

6

Logstash - Filters

```
filter {
    if (@[metadata][dest_index] == "kismet_device_activity") { 1
        date {
            match => ["kismet_device_base_last_time","UNIX"] target => "kismetdtg" 2
        }
        if [kismet_device_base_channel] == "FHSS" {
            prune { remove_field => ["[kismet_device_base_channel]"] } 3
            mutate { add_field => { "[kismet_device_base_channel_fhss]" => "FHSS" } }
            mutate { add_field => { "[kismet_device_base_channel]" => 0 } }
        }
        if [kismet_device_base_tags] == 0 {
            prune { remove_field => ["[kismet_device_base_tags]"] }
            mutate {
                add_field => { "[kismet_device_base_tags][notes]" => "NA" }
            }
        }
        prune {
            remove_field => ["[kismet_device_base_location][kismet_common_location_last]"]
            remove_field => ["[kismet_device_base_location][kismet_common_location_max_loc]"]
            remove_field => ["[kismet_device_base_location][kismet_common_location_min_loc]"] 4
        }
        mutate {
            add_field => {
                "my_fingerprint" => "devact_%{[kismet_device_base_key]}_%{[kismet_device_base_last_time]}" 5
            }
        }
    }
}
```

Logstash - Outputs

```
output {  
    elasticsearch {  
        hosts => ["localhost:9200"]  
        index => "%{@metadata}[dest_index]"  
        document_id => "%{[my_fingerprint]}"  
    }  
    #disable this once in prod stdout  
    #stdout { codec => rubydebug } # output any left overs  
}
```

Tips

- Read the (current) docs/SO/boards
- Trial and error
- Use repeatable test data
- Get paranoid about {}'s

Python

MUCH more customizable and powerful

Kismet and Elasticsearch libraries

Can check for and push indices

Can filter JSONs, conduct interim analysis, and create new fields

Can be scripted

Kibana Visualizations

Create an Index Pattern from an ES Index (Wildcards to match sub ES Indices)

Refresh field indexes when new fields appear

Creating a Visualization

- Types - Line/Bar, Heat Map, Coordinate Map
- Select Index Pattern
- Metrics (Y-Axis) - Count, Unique Count
- Buckets
 - X-Axis
 - Split Series
 - Split Chart
- Aggregations - Date Histogram, Terms
- Adjust “Metrics & Axis” & “Panel Settings”

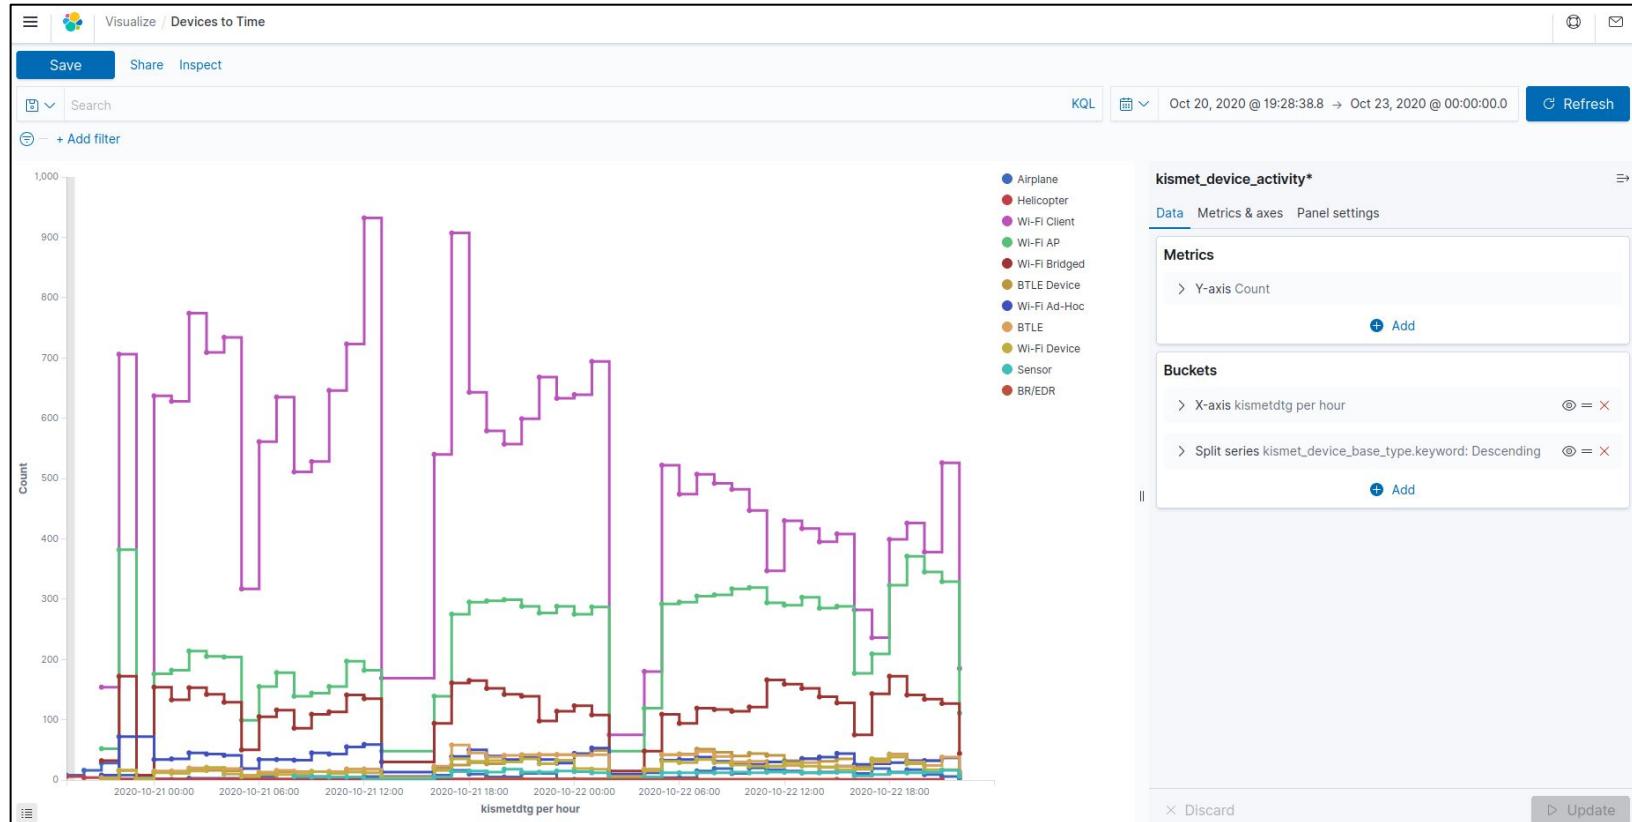
Key Fields

- kismet_device_base_key.keyword - Uniquely ID's a device
- kismet_device_base_type.keyword - Technology

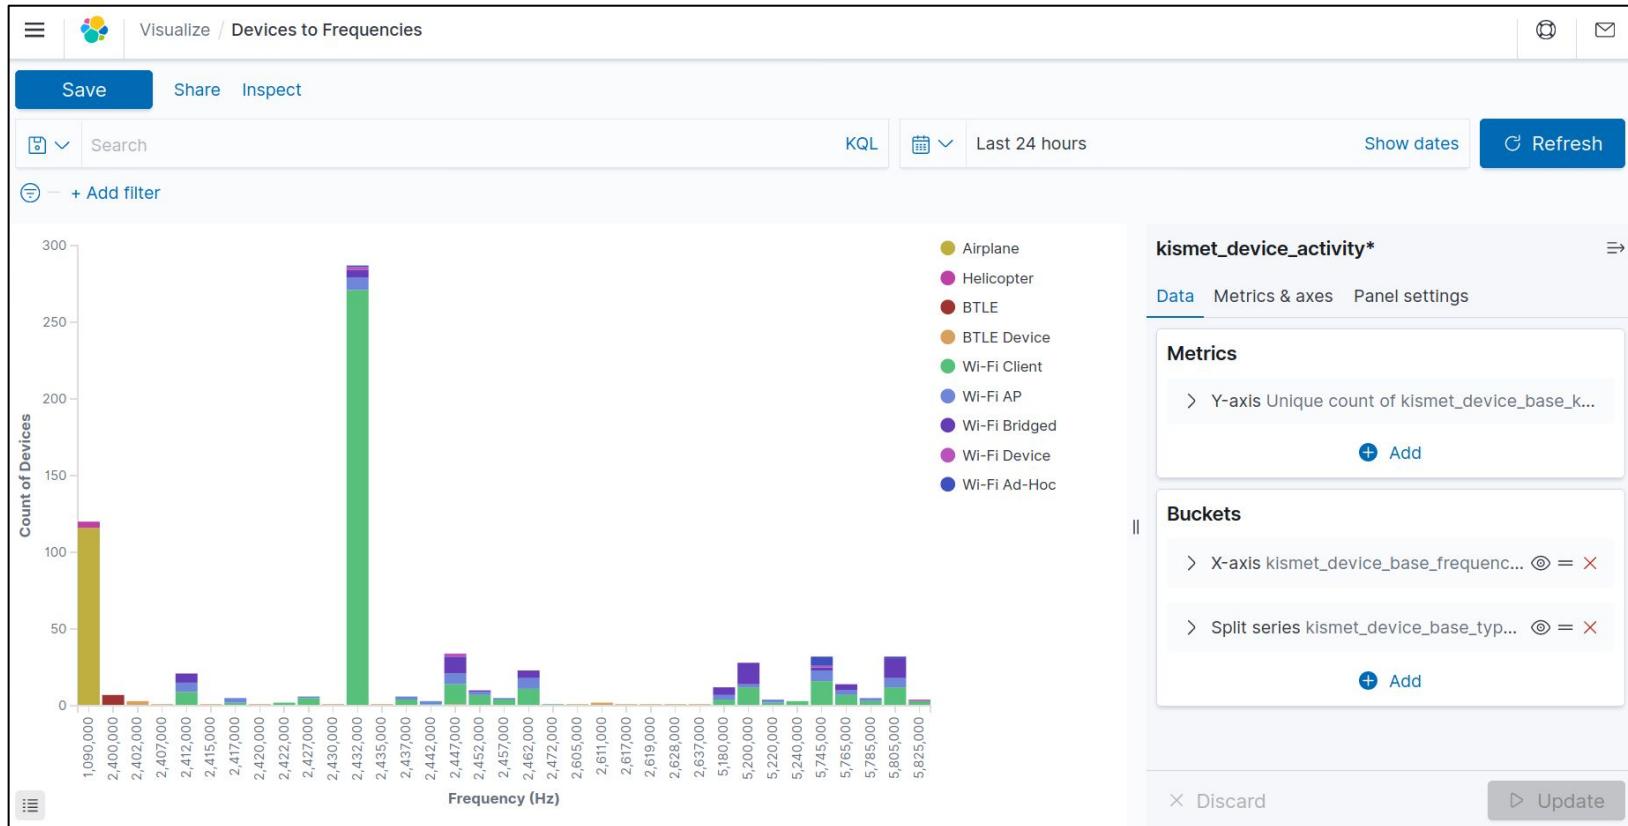
Tips

- Build in filters (but don't forget you did)
- Visualizations are tied to the Index Pattern's UUID...

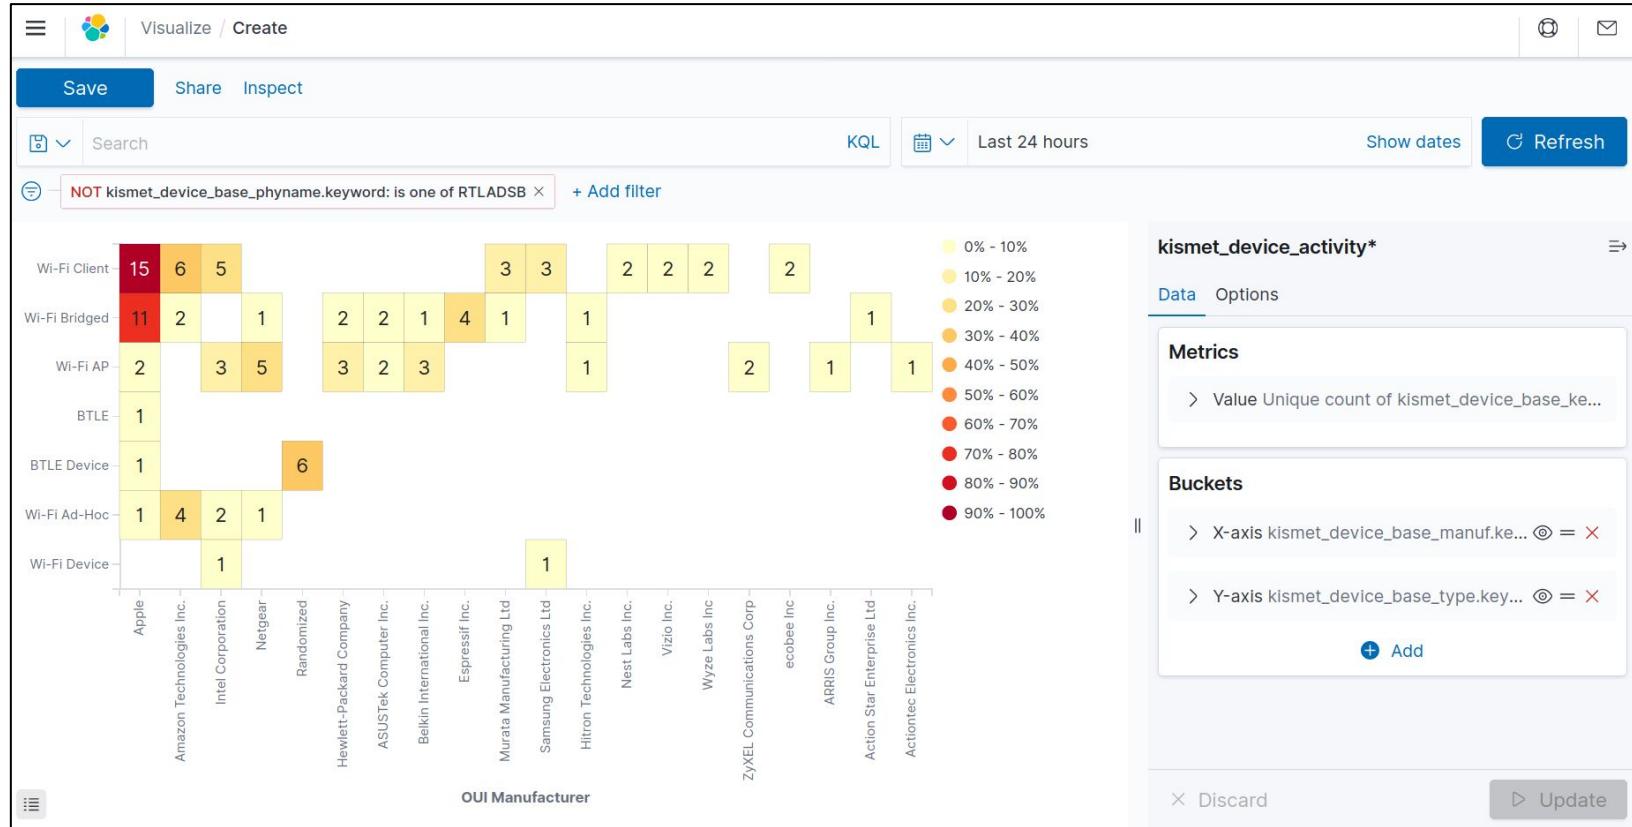
Line Graph - Device Types to Time



Bar Graph - Device Types to Frequency



Heat Map - Device Types to Manufacturer



Coordinate Map - Device Locations



Elasticsearch Organic WMS is limited to a zoom depth of 12

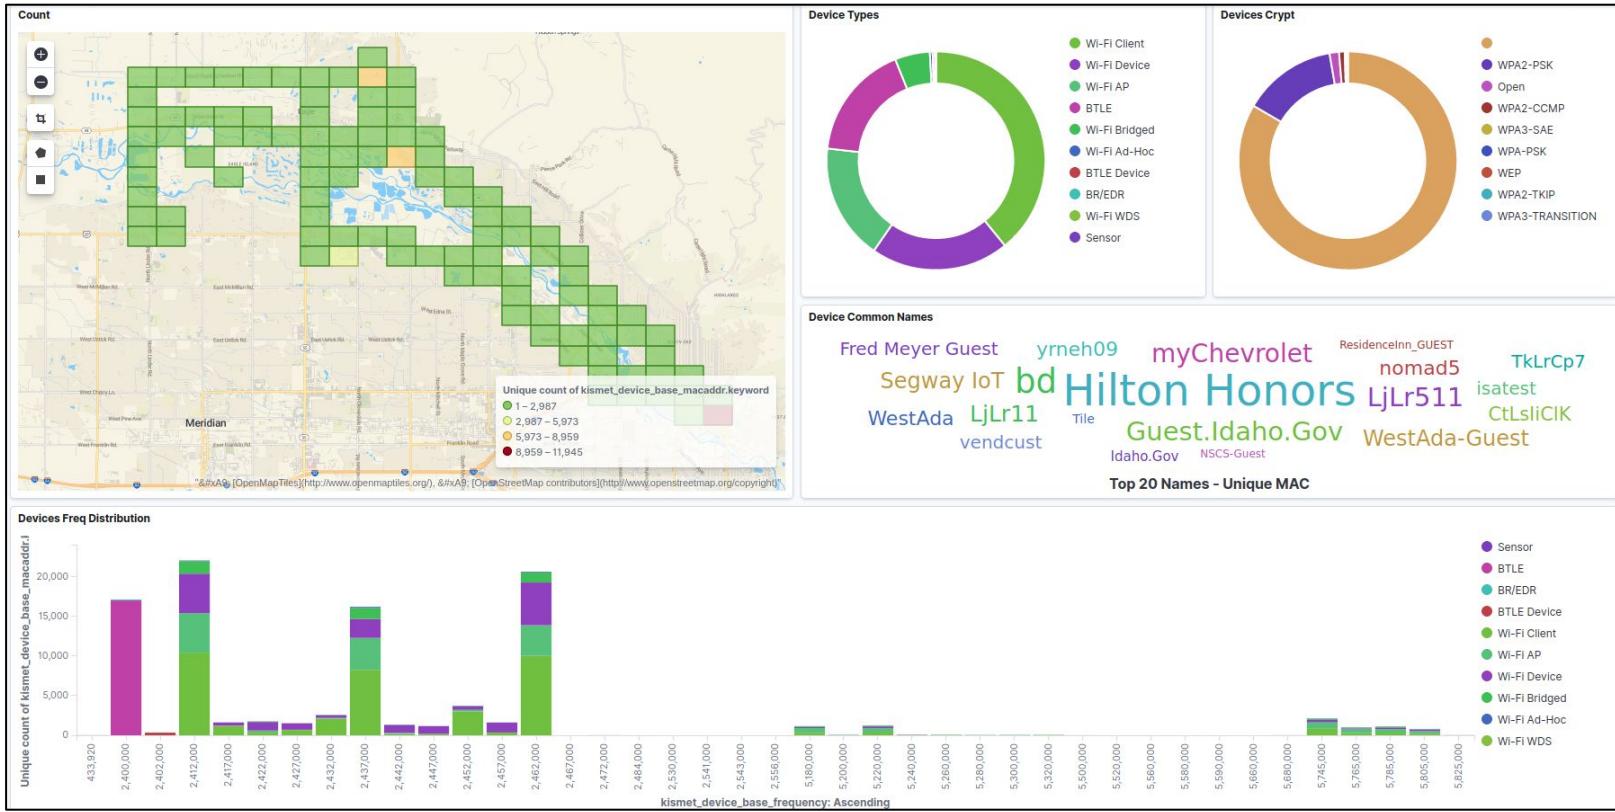
Custom WMS:

- <https://www.maptiler.com>
- WMS:
- Streets:
 - <https://api.maptiler.com/maps/streets/{z}/{x}/{y}.png?key=<your key>>
 - image/jpg
- Satellite:
 - WMS:
<https://api.maptiler.com/maps/hybrid/{z}/{x}/{y}.jpg?key=<your key>>
 - image/png

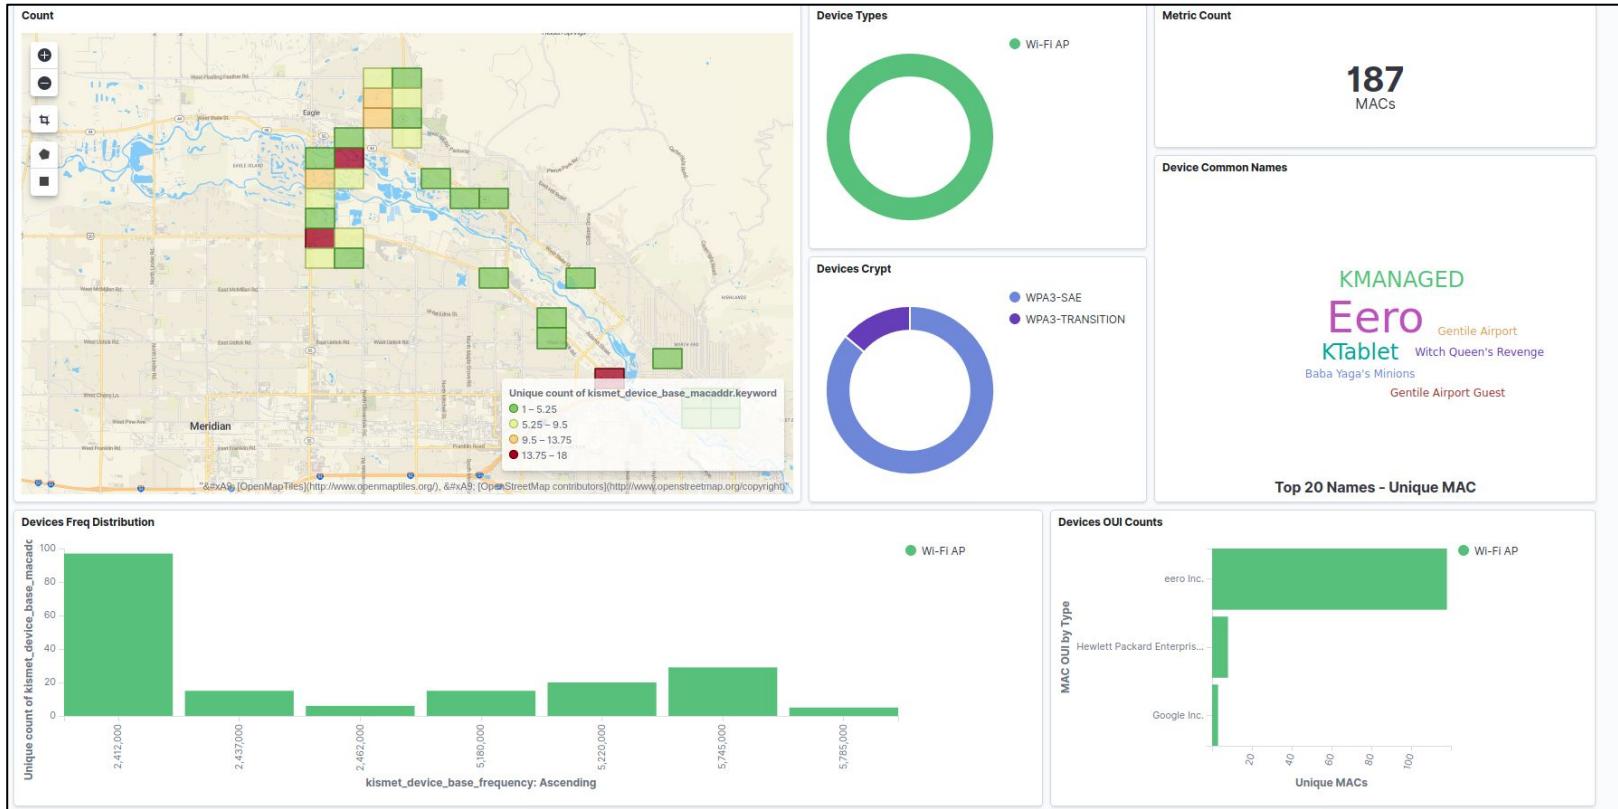
Dashboards - Home Monitoring



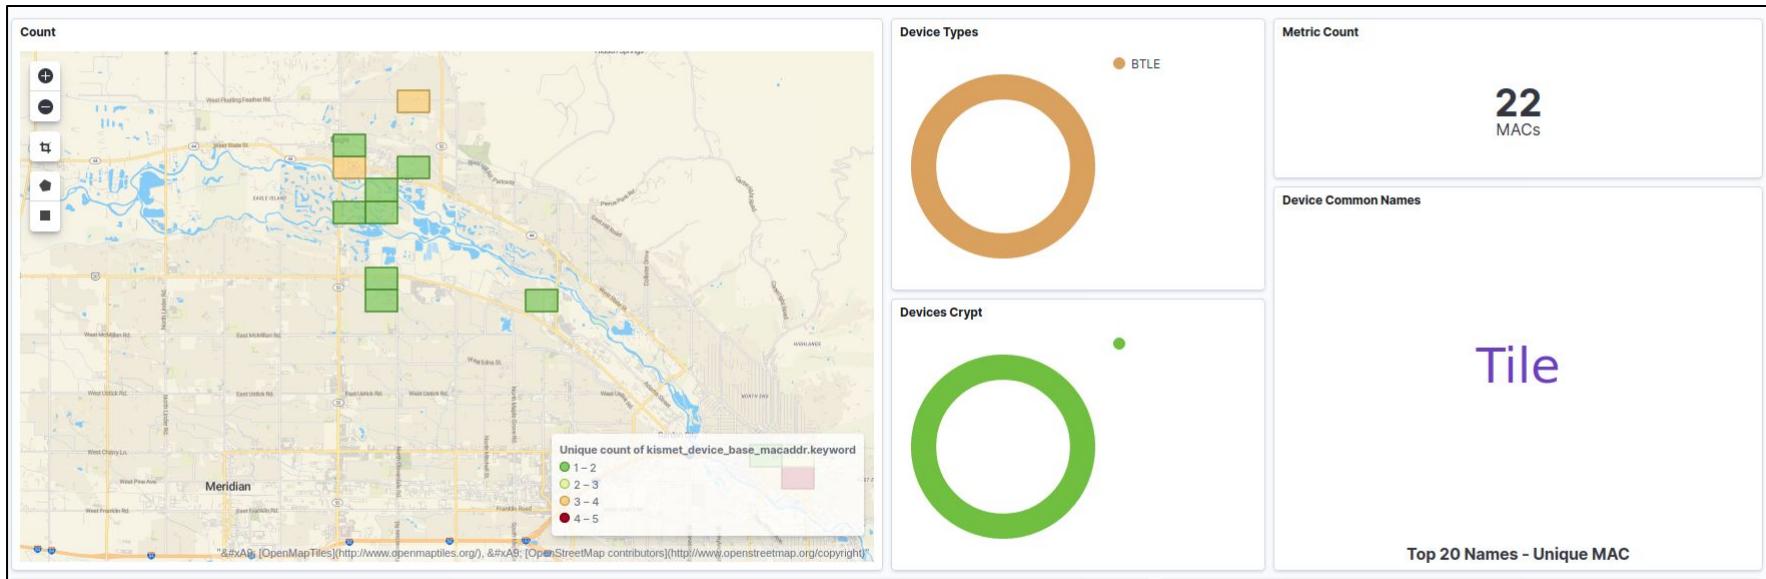
Dashboards - Spatial Trends



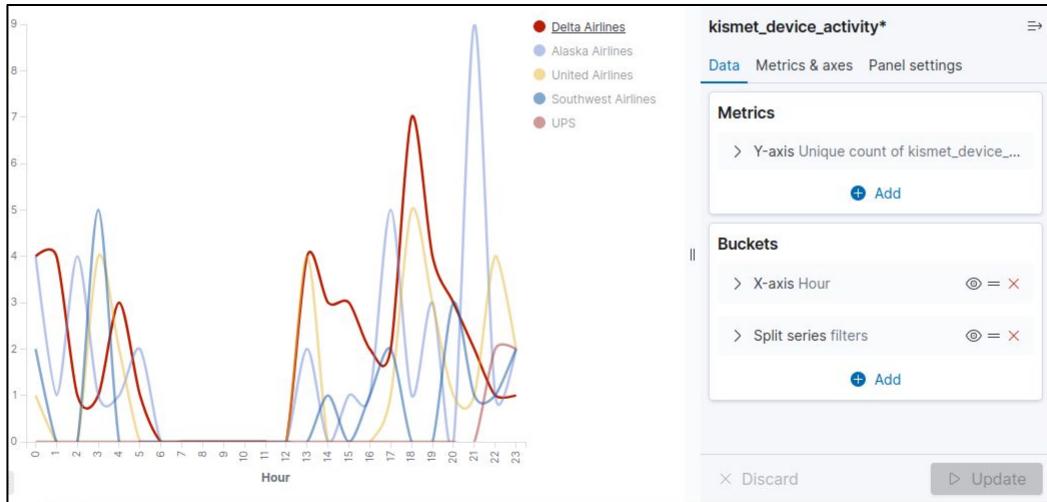
WPA3-SAE/Transition Distribution



Tile Devices



ADS-B: Airline Traffic Patterns



Scripted field in Kibana

- `doc['@timestamp'].value.getHour()`

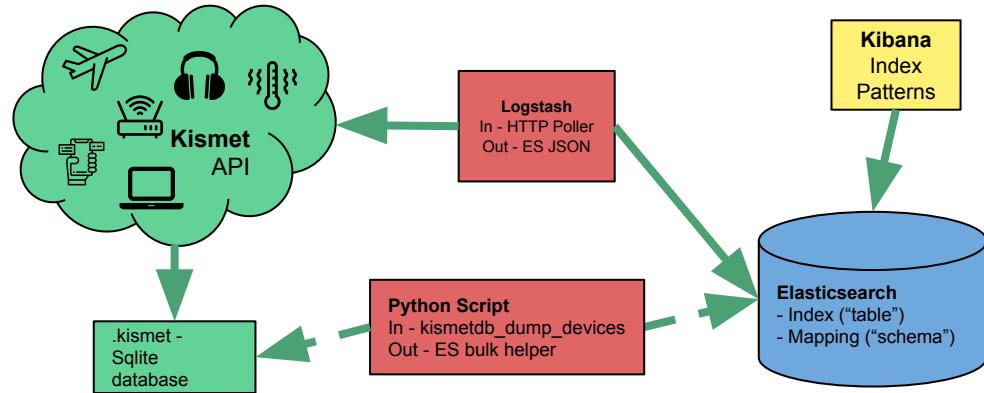
Filters

- Lucene not KQL
- `{"bool":{"should":[{"match_phrase": {"kismet_device_base_commonnme": "*DELTA*"}}]}}`

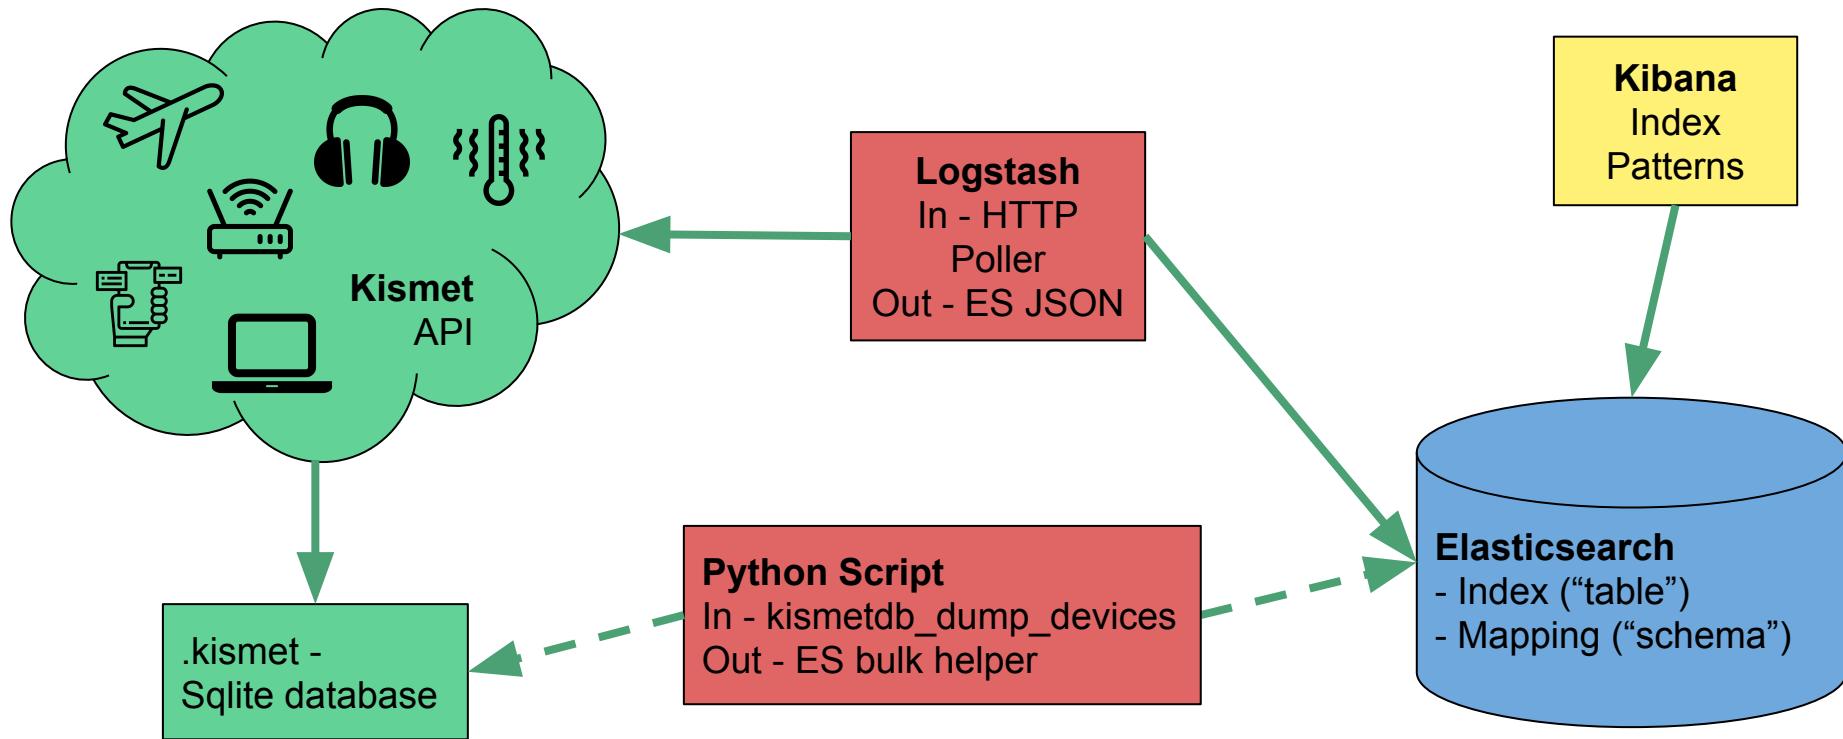
Order of Operations

1. Start Elasticsearch and Kibana
2. Create index and mapping in Elasticsearch
3. Configure Logstash pipeline
4. Start Kismet
5. Start Logstash
6. Once data is in ES, create a Kibana Index Pattern
7. Verify data in Kibana
8. Create visualizations and dashboards
9. Profit

WAWKELK Service Diagram



WAWKELK Service Diagram



Final Thoughts

Security - Careful how you expose/access this data

Github - Adding examples, lessons learned, and quick start guides

<https://github.com/developingchase/wawkelk>

Please contribute/add suggestions

Feel free to reach out with any questions:

- @developingchase
- Discord: chase#1342
- Via Github's Issue tracker

